This week I decided to add a local Ai assistant to my investment app. I don't have a plan for this in a production setting, so this week is purely for testing the idea. Its job is to explain things in plain English for users. If a user sees a term like "dividend yield" and doesn't understand its initial definition they can prompt the assistant to help them get a better understanding. The most important thing I learned about adding this AI was that not all answers would be beneficial for my users. I had to figure out how to make the assistant explain things simply for beginner users.

This idea came from my interviews from previous weeks. I kept thinking about how beginners were deterred by advanced wording and difficulty getting started. Now they can get more help without leaving the screen. I chose to run the model locally for testing since running it from the cloud would cost money. It is also beneficial because with a local model all questions are kept private on a personal device. This is important for users who are trusting an app for financial information.

Running a model locally was harder than I thought it would be. It is slower than a cloud model and gives less detailed responses since I am running a model on 16gb of VRAM. A problem I ran into was getting the AI to stay on topic. If a user asked a question out of scope of the app it would answer which it is not supposed to do. I learned quickly that AI implementation requires testing and hands on modification. In an important app such as financial the bot needs to stay on topic and give correct explanations, or it could easily lead a user down the wrong path.